

tktbmogzd

August 2, 2023

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: df=pd.read_csv("/content/drive/MyDrive/mydatasets/C1_ionosphere.csv")
df
```

```
[ ]:      1  0  0.99539 -0.05889  0.85243  0.02306  0.83398 -0.37708      1.1 \
0      1  0  1.00000 -0.18829  0.93035 -0.36156 -0.10868 -0.93597  1.00000
1      1  0  1.00000 -0.03365  1.00000  0.00485  1.00000 -0.12062  0.88965
2      1  0  1.00000 -0.45161  1.00000  1.00000  0.71216 -1.00000  0.00000
3      1  0  1.00000 -0.02401  0.94140  0.06531  0.92106 -0.23255  0.77152
4      1  0  0.02337 -0.00592 -0.09924 -0.11949 -0.00763 -0.11824  0.14706
..    ..  ..      ...      ...      ...      ...      ...      ...
345    1  0  0.83508  0.08298  0.73739 -0.14706  0.84349 -0.05567  0.90441
346    1  0  0.95113  0.00419  0.95183 -0.02723  0.93438 -0.01920  0.94590
347    1  0  0.94701 -0.00034  0.93207 -0.03227  0.95177 -0.03431  0.95584
348    1  0  0.90608 -0.01657  0.98122 -0.01989  0.95691 -0.03646  0.85746
349    1  0  0.84710  0.13533  0.73638 -0.06151  0.87873  0.08260  0.88928

      0.03760  ... -0.51171  0.41078 -0.46168  0.21266 -0.34090  0.42267 \
0    -0.04549  ... -0.26569 -0.20468 -0.18401 -0.19040 -0.11593 -0.16626
1      0.01198  ... -0.40220  0.58984 -0.22145  0.43100 -0.17365  0.60436
2      0.00000  ...  0.90695  0.51613  1.00000  1.00000 -0.20099  0.25682
3    -0.16399  ... -0.65158  0.13290 -0.53206  0.02431 -0.62197 -0.05707
4      0.06637  ... -0.01535 -0.03240  0.09223 -0.07859  0.00732  0.00000
..          ...  ...      ...      ...      ...      ...      ...
345 -0.04622  ... -0.04202  0.83479  0.00123  1.00000  0.12815  0.86660
346  0.01606  ...  0.01361  0.93522  0.04925  0.93159  0.08168  0.94066
```

```

347  0.02446 ...  0.03193  0.92489  0.02542  0.92120  0.02242  0.92459
348  0.00110 ... -0.02099  0.89147 -0.07760  0.82983 -0.17238  0.96022
349 -0.09139 ... -0.15114  0.81147 -0.04822  0.78207 -0.00703  0.75747

```

```

      -0.54487  0.18641 -0.45300 g
0      -0.06288 -0.13738 -0.02447 b
1      -0.24180  0.56045 -0.38238 g
2       1.00000 -0.32382  1.00000 b
3      -0.59573 -0.04608 -0.65697 g
4       0.00000 -0.00039  0.12011 b
..
345    -0.10714  0.90546 -0.04307 g
346    -0.00035  0.91483  0.04712 g
347     0.00442  0.92697 -0.00577 g
348    -0.03757  0.87403 -0.16243 g
349    -0.06678  0.85764 -0.06151 g

```

[350 rows x 35 columns]

```
[ ]: df.head()
```

```

[ ]:      1  0  0.99539 -0.05889  0.85243  0.02306  0.83398 -0.37708      1.1 \
0  1  0  1.00000 -0.18829  0.93035 -0.36156 -0.10868 -0.93597  1.00000
1  1  0  1.00000 -0.03365  1.00000  0.00485  1.00000 -0.12062  0.88965
2  1  0  1.00000 -0.45161  1.00000  1.00000  0.71216 -1.00000  0.00000
3  1  0  1.00000 -0.02401  0.94140  0.06531  0.92106 -0.23255  0.77152
4  1  0  0.02337 -0.00592 -0.09924 -0.11949 -0.00763 -0.11824  0.14706

```

```

      0.03760 ... -0.51171  0.41078 -0.46168  0.21266 -0.34090  0.42267 \
0 -0.04549 ... -0.26569 -0.20468 -0.18401 -0.19040 -0.11593 -0.16626
1  0.01198 ... -0.40220  0.58984 -0.22145  0.43100 -0.17365  0.60436
2  0.00000 ...  0.90695  0.51613  1.00000  1.00000 -0.20099  0.25682
3 -0.16399 ... -0.65158  0.13290 -0.53206  0.02431 -0.62197 -0.05707
4  0.06637 ... -0.01535 -0.03240  0.09223 -0.07859  0.00732  0.00000

```

```

      -0.54487  0.18641 -0.45300 g
0      -0.06288 -0.13738 -0.02447 b
1      -0.24180  0.56045 -0.38238 g
2       1.00000 -0.32382  1.00000 b
3      -0.59573 -0.04608 -0.65697 g
4       0.00000 -0.00039  0.12011 b

```

[5 rows x 35 columns]

1 Data Cleaning and Data Preprocessing

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 35 columns):
#   Column      Non-Null Count  Dtype
---  -
0   1            350 non-null    int64
1   0            350 non-null    int64
2   0.99539      350 non-null    float64
3   -0.05889     350 non-null    float64
4   0.85243      350 non-null    float64
5   0.02306      350 non-null    float64
6   0.83398      350 non-null    float64
7   -0.37708     350 non-null    float64
8   1.1          350 non-null    float64
9   0.03760      350 non-null    float64
10  0.85243.1    350 non-null    float64
11  -0.17755     350 non-null    float64
12  0.59755      350 non-null    float64
13  -0.44945     350 non-null    float64
14  0.60536      350 non-null    float64
15  -0.38223     350 non-null    float64
16  0.84356      350 non-null    float64
17  -0.38542     350 non-null    float64
18  0.58212      350 non-null    float64
19  -0.32192     350 non-null    float64
20  0.56971      350 non-null    float64
21  -0.29674     350 non-null    float64
22  0.36946      350 non-null    float64
23  -0.47357     350 non-null    float64
24  0.56811      350 non-null    float64
25  -0.51171     350 non-null    float64
26  0.41078      350 non-null    float64
27  -0.46168     350 non-null    float64
28  0.21266      350 non-null    float64
29  -0.34090     350 non-null    float64
30  0.42267      350 non-null    float64
31  -0.54487     350 non-null    float64
32  0.18641      350 non-null    float64
33  -0.45300     350 non-null    float64
34  g            350 non-null    object
dtypes: float64(32), int64(2), object(1)
memory usage: 95.8+ KB
```

```
[ ]: df.describe()
```

```
[ ]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	\
count	350.000000	350.0	350.000000	350.000000	350.000000	350.000000	
mean	0.891429	0.0	0.640330	0.044667	0.600350	0.116154	
std	0.311546	0.0	0.498059	0.442032	0.520431	0.461443	
min	0.000000	0.0	-1.000000	-1.000000	-1.000000	-1.000000	
25%	1.000000	0.0	0.471517	-0.065388	0.412555	-0.024868	
50%	1.000000	0.0	0.870795	0.016700	0.808620	0.021170	
75%	1.000000	0.0	1.000000	0.194727	1.000000	0.335317	
max	1.000000	0.0	1.000000	1.000000	1.000000	1.000000	

	0.83398	-0.37708	1.1	0.03760	...	0.56811	\
count	350.000000	350.000000	350.000000	350.000000	...	350.000000	
mean	0.549284	0.120779	0.510453	0.181756	...	0.395643	
std	0.493124	0.520816	0.507117	0.484482	...	0.579206	
min	-1.000000	-1.000000	-1.000000	-1.000000	...	-1.000000	
25%	0.209105	-0.053483	0.086785	-0.049003	...	0.000000	
50%	0.728000	0.015085	0.682430	0.017550	...	0.549175	
75%	0.970445	0.451572	0.950555	0.536192	...	0.907165	
max	1.000000	1.000000	1.000000	1.000000	...	1.000000	

	-0.51171	0.41078	-0.46168	0.21266	-0.34090	0.42267	\
count	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000	
mean	-0.069928	0.542015	-0.068417	0.378919	-0.027013	0.352313	
std	0.508675	0.516896	0.550411	0.576642	0.508425	0.572289	
min	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	
25%	-0.323745	0.283612	-0.428992	0.000000	-0.234935	0.000000	
50%	-0.014915	0.708530	-0.017685	0.499215	0.000000	0.446875	
75%	0.157922	0.999972	0.154862	0.884572	0.154218	0.859490	
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	

	-0.54487	0.18641	-0.45300
count	350.000000	350.000000	350.000000
mean	-0.002248	0.349829	0.015816
std	0.513491	0.523339	0.468338
min	-1.000000	-1.000000	-1.000000
25%	-0.239347	0.000000	-0.161013
50%	0.000000	0.413115	0.000000
75%	0.200935	0.816778	0.172105
max	1.000000	1.000000	1.000000

[8 rows x 34 columns]

```
[ ]: df.columns
```

```
[ ]: Index(['1', '0', '0.99539', '-0.05889', '0.85243', '0.02306', '0.83398',
          '-0.37708', '1.1', '0.03760', '0.85243.1', '-0.17755', '0.59755',
          '-0.44945', '0.60536', '-0.38223', '0.84356', '-0.38542', '0.58212',
          '-0.32192', '0.56971', '-0.29674', '0.36946', '-0.47357', '0.56811',
          '-0.51171', '0.41078', '-0.46168', '0.21266', '-0.34090', '0.42267',
          '-0.54487', '0.18641', '-0.45300', 'g'],
          dtype='object')
```

```
[ ]: feature_matrix = df.iloc[:,0:34]
      target_vector = df.iloc[:,-1]
```

```
[ ]: fs = StandardScaler().fit_transform(feature_matrix)
      logr = LogisticRegression()
      logr.fit(fs,target_vector)
```

```
[ ]: LogisticRegression()
```

```
[ ]: observation=[[1.0,0.0,1.0,-0.18829,0.93035,
                  -0.36156,
                  -0.10868,
                  -0.93597,
                  1.0,
                  -0.04549,
                  0.50874,
                  -0.67743,
                  0.34432,
                  -0.69707,
                  -0.51685,
                  -0.97515,
                  0.05499,
                  -0.62237,
                  0.33109,
                  -1.0,
                  -0.13151,
                  -0.453,
                  -0.18056,
                  -0.35734,
                  -0.20332,
                  -0.26569,
                  -0.20468,
                  -0.18401,
                  -0.1904,
                  -0.11593,
                  -0.16626,
                  -0.06288,
                  -0.13738,
                  -0.02447]]
```

```
prediction = logr.predict(observation)
print(prediction)
```

```
['g']
```

```
[ ]: logr.classes_
```

```
[ ]: array(['b', 'g'], dtype=object)
```

```
[ ]: logr.predict_proba(observation)
```

```
[ ]: array([[0.07006552, 0.92993448]])
```

Random Forest

```
[ ]: df['g'].value_counts()
```

```
[ ]: g    224
      b    126
      Name: g, dtype: int64
```

```
[ ]: x=df.drop('g', axis=1)
      y=df['g']
```

```
[ ]: g1={"g":{"g":1, "b":2}}
      df=df.replace(g1)
      df
```

```
[ ]:      1  0  0.99539  -0.05889  0.85243  0.02306  0.83398  -0.37708      1.1  \
0      1  0  1.00000  -0.18829  0.93035  -0.36156  -0.10868  -0.93597  1.00000
1      1  0  1.00000  -0.03365  1.00000  0.00485  1.00000  -0.12062  0.88965
2      1  0  1.00000  -0.45161  1.00000  1.00000  0.71216  -1.00000  0.00000
3      1  0  1.00000  -0.02401  0.94140  0.06531  0.92106  -0.23255  0.77152
4      1  0  0.02337  -0.00592  -0.09924  -0.11949  -0.00763  -0.11824  0.14706
..  ..  ..      ...      ...      ...      ...      ...      ...
345    1  0  0.83508   0.08298  0.73739  -0.14706  0.84349  -0.05567  0.90441
346    1  0  0.95113   0.00419  0.95183  -0.02723  0.93438  -0.01920  0.94590
347    1  0  0.94701  -0.00034  0.93207  -0.03227  0.95177  -0.03431  0.95584
348    1  0  0.90608  -0.01657  0.98122  -0.01989  0.95691  -0.03646  0.85746
349    1  0  0.84710   0.13533  0.73638  -0.06151  0.87873   0.08260  0.88928

      0.03760  ...  -0.51171  0.41078  -0.46168  0.21266  -0.34090  0.42267  \
0  -0.04549  ...  -0.26569  -0.20468  -0.18401  -0.19040  -0.11593  -0.16626
1   0.01198  ...  -0.40220  0.58984  -0.22145  0.43100  -0.17365  0.60436
2   0.00000  ...   0.90695  0.51613   1.00000  1.00000  -0.20099  0.25682
3  -0.16399  ...  -0.65158  0.13290  -0.53206  0.02431  -0.62197  -0.05707
4   0.06637  ...  -0.01535  -0.03240   0.09223  -0.07859   0.00732  0.00000
..      ...  ...      ...      ...      ...      ...      ...      ...
```

```

345 -0.04622 ... -0.04202 0.83479 0.00123 1.00000 0.12815 0.86660
346 0.01606 ... 0.01361 0.93522 0.04925 0.93159 0.08168 0.94066
347 0.02446 ... 0.03193 0.92489 0.02542 0.92120 0.02242 0.92459
348 0.00110 ... -0.02099 0.89147 -0.07760 0.82983 -0.17238 0.96022
349 -0.09139 ... -0.15114 0.81147 -0.04822 0.78207 -0.00703 0.75747

```

```

      -0.54487 0.18641 -0.45300 g
0 -0.06288 -0.13738 -0.02447 2
1 -0.24180 0.56045 -0.38238 1
2 1.00000 -0.32382 1.00000 2
3 -0.59573 -0.04608 -0.65697 1
4 0.00000 -0.00039 0.12011 2
.. ...
345 -0.10714 0.90546 -0.04307 1
346 -0.00035 0.91483 0.04712 1
347 0.00442 0.92697 -0.00577 1
348 -0.03757 0.87403 -0.16243 1
349 -0.06678 0.85764 -0.06151 1

```

[350 rows x 35 columns]

```
[ ]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
[ ]: RandomForestClassifier()
```

```
[ ]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                  'n_estimators': [10,20,30,40,50]
                  }
```

```
[ ]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
[ ]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                param_grid={'max_depth': [1, 2, 3, 4, 5],
                            'min_samples_leaf': [5, 10, 15, 20, 25],
                            'n_estimators': [10, 20, 30, 40, 50]},
                scoring='accuracy')
```

```
[ ]: grid_search.best_score_
```

```
[ ]: 0.9344262295081966
```

```
[ ]: rfc_best = grid_search.best_estimator_
```

```
[ ]: from sklearn.tree import plot_tree
plt.figure(figsize=(89,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes', 'No'], filled=True)
```

```
[ ]: [Text(0.42105263157894735, 0.9166666666666666, '0.42267 <= 0.12\ngini =
0.474\nsamples = 157\nvalue = [94, 150]\nclass = No'),
Text(0.21052631578947367, 0.75, '0.42267 <= -0.018\ngini = 0.39\nsamples =
50\nvalue = [58, 21]\nclass = Yes'),
Text(0.10526315789473684, 0.5833333333333333, '0.83398 <= 0.346\ngini =
0.5\nsamples = 29\nvalue = [19, 20]\nclass = No'),
Text(0.05263157894736842, 0.4166666666666667, 'gini = 0.26\nsamples = 8\nvalue
= [11, 2]\nclass = Yes'),
Text(0.15789473684210525, 0.4166666666666667, '0.42267 <= -0.732\ngini =
0.426\nsamples = 21\nvalue = [8, 18]\nclass = No'),
Text(0.10526315789473684, 0.25, 'gini = 0.32\nsamples = 7\nvalue = [8,
2]\nclass = Yes'),
Text(0.21052631578947367, 0.25, 'gini = 0.0\nsamples = 14\nvalue = [0,
16]\nclass = No'),
Text(0.3157894736842105, 0.5833333333333333, '0.60536 <= -0.018\ngini =
0.049\nsamples = 21\nvalue = [39, 1]\nclass = Yes'),
Text(0.2631578947368421, 0.4166666666666667, 'gini = 0.219\nsamples = 5\nvalue
= [7, 1]\nclass = Yes'),
Text(0.3684210526315789, 0.4166666666666667, 'gini = 0.0\nsamples = 16\nvalue =
[32, 0]\nclass = Yes'),
Text(0.631578947368421, 0.75, '-0.34090 <= -0.934\ngini = 0.341\nsamples =
107\nvalue = [36, 129]\nclass = No'),
Text(0.5263157894736842, 0.5833333333333333, '-0.44945 <= -0.932\ngini =
0.278\nsamples = 11\nvalue = [15, 3]\nclass = Yes'),
Text(0.47368421052631576, 0.4166666666666667, 'gini = 0.0\nsamples = 5\nvalue =
[7, 0]\nclass = Yes'),
Text(0.5789473684210527, 0.4166666666666667, 'gini = 0.397\nsamples = 6\nvalue
= [8, 3]\nclass = Yes'),
Text(0.7368421052631579, 0.5833333333333333, '0.99539 <= 0.133\ngini =
0.245\nsamples = 96\nvalue = [21, 126]\nclass = No'),
Text(0.6842105263157895, 0.4166666666666667, 'gini = 0.0\nsamples = 5\nvalue =
[8, 0]\nclass = Yes'),
Text(0.7894736842105263, 0.4166666666666667, '0.85243 <= 0.995\ngini =
0.17\nsamples = 91\nvalue = [13, 126]\nclass = No'),
Text(0.6842105263157895, 0.25, '0.42267 <= 0.433\ngini = 0.042\nsamples =
61\nvalue = [2, 92]\nclass = No'),
Text(0.631578947368421, 0.08333333333333333, 'gini = 0.147\nsamples = 19\nvalue
= [2, 23]\nclass = No'),
```



```

Text(0.7368421052631579, 0.08333333333333333, 'gini = 0.0\nsamples = 42\nvalue = [0, 69]\nclass = No'),
Text(0.8947368421052632, 0.25, '-0.05889 <= 0.348\ngini = 0.369\nsamples = 30\nvalue = [11, 34]\nclass = No'),
Text(0.8421052631578947, 0.08333333333333333, 'gini = 0.234\nsamples = 24\nvalue = [5, 32]\nclass = No'),
Text(0.9473684210526315, 0.08333333333333333, 'gini = 0.375\nsamples = 6\nvalue = [6, 2]\nclass = Yes')]

```

