

kzzhuuzhp

August 2, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: df=pd.read_csv("/content/drive/MyDrive/mydatasets/C4_framingham.csv")
df
```

```
[3]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
0	1	39	4.0	0	0.0	0.0	
1	0	46	2.0	0	0.0	0.0	
2	1	48	1.0	1	20.0	0.0	
3	0	61	3.0	1	30.0	0.0	
4	0	46	3.0	1	23.0	0.0	
...	
4233	1	50	1.0	1	1.0	0.0	
4234	1	51	3.0	1	43.0	0.0	
4235	0	48	2.0	1	20.0	NaN	
4236	0	44	1.0	1	15.0	0.0	
4237	0	52	2.0	0	0.0	0.0	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...	
4233	0	1	0	313.0	179.0	92.0	25.97	
4234	0	0	0	207.0	126.5	80.0	19.71	
4235	0	0	0	248.0	131.0	72.0	22.00	

4236	0	0	0	210.0	126.5	87.0	19.16
4237	0	0	0	269.0	133.5	83.0	21.47

	heartRate	glucose	TenYearCHD
0	80.0	77.0	0
1	95.0	76.0	0
2	75.0	70.0	0
3	65.0	103.0	1
4	85.0	85.0	0
...
4233	66.0	86.0	1
4234	65.0	68.0	0
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0

[4238 rows x 16 columns]

```
[4]: df.head()
```

```
[4]:   male  age  education  currentSmoker  cigsPerDay  BPMeds  prevalentStroke  \
0     1   39         4.0              0          0.0     0.0              0
1     0   46         2.0              0          0.0     0.0              0
2     1   48         1.0              1         20.0     0.0              0
3     0   61         3.0              1         30.0     0.0              0
4     0   46         3.0              1         23.0     0.0              0
```

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

	TenYearCHD
0	0
1	0
2	0
3	1
4	0

1 Data Cleaning and Data Preprocessing

```
[5]: df.dropna(inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3656 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  3656 non-null   int64
1   age                   3656 non-null   int64
2   education             3656 non-null   float64
3   currentSmoker         3656 non-null   int64
4   cigsPerDay            3656 non-null   float64
5   BPMeds                3656 non-null   float64
6   prevalentStroke       3656 non-null   int64
7   prevalentHyp          3656 non-null   int64
8   diabetes              3656 non-null   int64
9   totChol               3656 non-null   float64
10  sysBP                 3656 non-null   float64
11  diaBP                 3656 non-null   float64
12  BMI                   3656 non-null   float64
13  heartRate             3656 non-null   float64
14  glucose               3656 non-null   float64
15  TenYearCHD            3656 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 485.6 KB
```

```
[6]: df.describe()
```

```
[6]:
```

	male	age	education	currentSmoker	cigsPerDay	\
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	
mean	0.443654	49.557440	1.979759	0.489059	9.022155	
std	0.496883	8.561133	1.022657	0.499949	11.918869	
min	0.000000	32.000000	1.000000	0.000000	0.000000	
25%	0.000000	42.000000	1.000000	0.000000	0.000000	
50%	0.000000	49.000000	2.000000	0.000000	0.000000	
75%	1.000000	56.000000	3.000000	1.000000	20.000000	
max	1.000000	70.000000	4.000000	1.000000	70.000000	

	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	\
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	
mean	0.030361	0.005744	0.311543	0.027079	236.873085	
std	0.171602	0.075581	0.463187	0.162335	44.096223	
min	0.000000	0.000000	0.000000	0.000000	113.000000	
25%	0.000000	0.000000	0.000000	0.000000	206.000000	

50%	0.000000	0.000000	0.000000	0.000000	234.000000
75%	0.000000	0.000000	1.000000	0.000000	263.250000
max	1.000000	1.000000	1.000000	1.000000	600.000000

	sysBP	diaBP	BMI	heartRate	glucose \
count	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000
mean	132.368025	82.912062	25.784185	75.730580	81.856127
std	22.092444	11.974825	4.065913	11.982952	23.910128
min	83.500000	48.000000	15.540000	44.000000	40.000000
25%	117.000000	75.000000	23.080000	68.000000	71.000000
50%	128.000000	82.000000	25.380000	75.000000	78.000000
75%	144.000000	90.000000	28.040000	82.000000	87.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	3656.000000
mean	0.152352
std	0.359411
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[7]: df.columns
```

```
[7]: Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds',
        'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP',
        'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'],
        dtype='object')
```

```
[8]: feature_matrix = df.iloc[:,0:15]
      target_vector = df.iloc[:,-1]
```

```
[9]: fs = StandardScaler().fit_transform(feature_matrix)
      logr = LogisticRegression()
      logr.fit(fs,target_vector)
```

```
[9]: LogisticRegression()
```

```
[10]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]]
      prediction = logr.predict(observation)
      print(prediction)
```

```
[1]
```

```
[11]: logr.classes_
```

```
[11]: array([0, 1])
```

```
[12]: logit.predict_proba(observation)
```

```
[12]: array([[2.21478351e-04, 9.99778522e-01]])
```

Random Forest

```
[18]: x = df.iloc[:,0:15]
      y = df.iloc[:,-1]
```

```
[19]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
[20]: from sklearn.ensemble import RandomForestClassifier
      rfc = RandomForestClassifier()
      rfc.fit(x_train,y_train)
```

```
[20]: RandomForestClassifier()
```

```
[21]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                  'n_estimators': [10,20,30,40,50]
                  }
```

```
[22]: from sklearn.model_selection import GridSearchCV
      grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
      grid_search.fit(x_train,y_train)
```

```
[22]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [1, 2, 3, 4, 5],
                              'min_samples_leaf': [5, 10, 15, 20, 25],
                              'n_estimators': [10, 20, 30, 40, 50]},
                  scoring='accuracy')
```

```
[23]: grid_search.best_score_
```

```
[23]: 0.8499410550234558
```

```
[24]: rfc_best = grid_search.best_estimator_
```

```
[25]: from sklearn.tree import plot_tree
      plt.figure(figsize=(89,40))
      plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes', 'No'],
                filled=True)
```

```

[25]: [Text(0.5625, 0.9166666666666666, 'glucose <= 146.0\ngini = 0.259\nsamples =
1637\nvalue = [2167, 392]\nclass = Yes'),
Text(0.5267857142857143, 0.75, 'prevalentHyp <= 0.5\ngini = 0.25\nsamples =
1616\nvalue = [2157, 370]\nclass = Yes'),
Text(0.2857142857142857, 0.5833333333333334, 'age <= 48.5\ngini =
0.186\nsamples = 1110\nvalue = [1545, 179]\nclass = Yes'),
Text(0.14285714285714285, 0.4166666666666667, 'cigsPerDay <= 9.5\ngini =
0.113\nsamples = 660\nvalue = [955, 61]\nclass = Yes'),
Text(0.07142857142857142, 0.25, 'glucose <= 91.5\ngini = 0.071\nsamples =
356\nvalue = [521, 20]\nclass = Yes'),
Text(0.03571428571428571, 0.08333333333333333, 'gini = 0.052\nsamples =
323\nvalue = [473, 13]\nclass = Yes'),
Text(0.10714285714285714, 0.08333333333333333, 'gini = 0.222\nsamples =
33\nvalue = [48, 7]\nclass = Yes'),
Text(0.21428571428571427, 0.25, 'heartRate <= 90.5\ngini = 0.158\nsamples =
304\nvalue = [434, 41]\nclass = Yes'),
Text(0.17857142857142858, 0.08333333333333333, 'gini = 0.168\nsamples =
284\nvalue = [403, 41]\nclass = Yes'),
Text(0.25, 0.08333333333333333, 'gini = 0.0\nsamples = 20\nvalue = [31,
0]\nclass = Yes'),
Text(0.42857142857142855, 0.4166666666666667, 'cigsPerDay <= 2.5\ngini =
0.278\nsamples = 450\nvalue = [590, 118]\nclass = Yes'),
Text(0.35714285714285715, 0.25, 'sysBP <= 145.5\ngini = 0.2\nsamples =
260\nvalue = [369, 47]\nclass = Yes'),
Text(0.32142857142857145, 0.08333333333333333, 'gini = 0.183\nsamples =
245\nvalue = [352, 40]\nclass = Yes'),
Text(0.39285714285714285, 0.08333333333333333, 'gini = 0.413\nsamples =
15\nvalue = [17, 7]\nclass = Yes'),
Text(0.5, 0.25, 'glucose <= 63.5\ngini = 0.368\nsamples = 190\nvalue = [221,
71]\nclass = Yes'),
Text(0.4642857142857143, 0.08333333333333333, 'gini = 0.0\nsamples = 19\nvalue
= [28, 0]\nclass = Yes'),
Text(0.5357142857142857, 0.08333333333333333, 'gini = 0.393\nsamples =
171\nvalue = [193, 71]\nclass = Yes'),
Text(0.7678571428571429, 0.5833333333333334, 'age <= 54.5\ngini =
0.363\nsamples = 506\nvalue = [612, 191]\nclass = Yes'),
Text(0.6785714285714286, 0.4166666666666667, 'sysBP <= 176.75\ngini =
0.259\nsamples = 269\nvalue = [360, 65]\nclass = Yes'),
Text(0.6428571428571429, 0.25, 'male <= 0.5\ngini = 0.222\nsamples = 250\nvalue
= [343, 50]\nclass = Yes'),
Text(0.6071428571428571, 0.08333333333333333, 'gini = 0.123\nsamples =
123\nvalue = [184, 13]\nclass = Yes'),
Text(0.6785714285714286, 0.08333333333333333, 'gini = 0.306\nsamples =
127\nvalue = [159, 37]\nclass = Yes'),
Text(0.7142857142857143, 0.25, 'gini = 0.498\nsamples = 19\nvalue = [17,
15]\nclass = Yes'),
Text(0.8571428571428571, 0.4166666666666667, 'glucose <= 69.5\ngini =

```

```

0.444\nsamples = 237\nnvalue = [252, 126]\nclass = Yes'),
Text(0.7857142857142857, 0.25, 'heartRate <= 75.5\ngini = 0.124\nsamples =
28\nnvalue = [42, 3]\nclass = Yes'),
Text(0.75, 0.08333333333333333, 'gini = 0.153\nsamples = 16\nnvalue = [22,
2]\nclass = Yes'),
Text(0.8214285714285714, 0.08333333333333333, 'gini = 0.091\nsamples =
12\nnvalue = [20, 1]\nclass = Yes'),
Text(0.9285714285714286, 0.25, 'totChol <= 318.5\ngini = 0.466\nsamples =
209\nnvalue = [210, 123]\nclass = Yes'),
Text(0.8928571428571429, 0.08333333333333333, 'gini = 0.452\nsamples =
191\nnvalue = [198, 104]\nclass = Yes'),
Text(0.9642857142857143, 0.08333333333333333, 'gini = 0.475\nsamples =
18\nnvalue = [12, 19]\nclass = No'),
Text(0.5982142857142857, 0.75, 'gini = 0.43\nsamples = 21\nnvalue = [10,
22]\nclass = No')]

```

