

hd4xzrz5t

August 2, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: df=pd.read_csv("/content/drive/MyDrive/mydatasets/C5_health care diabetes.csv")
df
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0

```

766          0.349  47      1
767          0.315  23      0

```

[768 rows x 9 columns]

```
[4]: df.head()
```

```

[4]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0           6      148            72           35         0  33.6
1           1       85            66           29         0  26.6
2           8      183            64            0         0  23.3
3           1       89            66           23        94  28.1
4           0      137            40           35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627    50         1
1                0.351    31         0
2                0.672    32         1
3                0.167    21         0
4                2.288    33         1

```

1 Data Cleaning and Data Preprocessing

```
[5]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                  768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                  768 non-null   int64
8   Outcome              768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[6]: df.describe()
```

```

[6]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  \
count  768.000000  768.000000    768.000000    768.000000  768.000000

```

mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[7]: df.columns
```

```
[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
        'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
        dtype='object')
```

```
[8]: feature_matrix = df.iloc[:,0:8]
      target_vector = df.iloc[:,-1]
```

```
[9]: fs = StandardScaler().fit_transform(feature_matrix)
      logr = LogisticRegression()
      logr.fit(fs,target_vector)
```

```
[9]: LogisticRegression()
```

```
[10]: observation=[[1,2,3,4,5,6,7,8]]
      prediction = logr.predict(observation)
      print(prediction)
```

```
[1]
```

```
[11]: logr.classes_
```

```
[11]: array([0, 1])
```

```
[12]: logr.predict_proba(observation)
```

```
[12]: array([[2.92369487e-04, 9.99707631e-01]])
```

Random Forest

```
[13]: x = df.iloc[:,0:8]
      y = df.iloc[:,-1]
```

```
[14]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
[15]: from sklearn.ensemble import RandomForestClassifier
      rfc = RandomForestClassifier()
      rfc.fit(x_train,y_train)
```

```
[15]: RandomForestClassifier()
```

```
[16]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators': [10,20,30,40,50]
                    }
```

```
[17]: from sklearn.model_selection import GridSearchCV
      grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
      grid_search.fit(x_train,y_train)
```

```
[17]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [1, 2, 3, 4, 5],
                              'min_samples_leaf': [5, 10, 15, 20, 25],
                              'n_estimators': [10, 20, 30, 40, 50]},
                  scoring='accuracy')
```

```
[18]: grid_search.best_score_
```

```
[18]: 0.7765424735060755
```

```
[19]: rfc_best = grid_search.best_estimator_
```

```
[20]: from sklearn.tree import plot_tree
      plt.figure(figsize=(89,40))
      plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes', 'No'], filled=True)
```

```
[20]: [Text(0.49107142857142855, 0.9, 'Age <= 28.5\ngini = 0.456\nsamples = 342\nvalue = [348, 189]\nclass = Yes'),
      Text(0.26785714285714285, 0.7, 'Glucose <= 141.5\ngini = 0.321\nsamples = 160\nvalue = [211, 53]\nclass = Yes'),
      Text(0.14285714285714285, 0.5, 'SkinThickness <= 5.0\ngini = 0.197\nsamples = 137\nvalue = [201, 25]\nclass = Yes'),
      Text(0.07142857142857142, 0.3, 'BloodPressure <= 79.0\ngini = 0.363\nsamples = 25\nvalue = [32, 10]\nclass = Yes'),
      Text(0.03571428571428571, 0.1, 'gini = 0.307\nsamples = 20\nvalue = [30,
```

```

7]\nclass = Yes'),
  Text(0.10714285714285714, 0.1, 'gini = 0.48\nsamples = 5\nvalue = [2, 3]\nclass
= No'),
  Text(0.21428571428571427, 0.3, 'DiabetesPedigreeFunction <= 0.67\ngini =
0.15\nsamples = 112\nvalue = [169, 15]\nclass = Yes'),
  Text(0.17857142857142858, 0.1, 'gini = 0.109\nsamples = 91\nvalue = [147,
9]\nclass = Yes'),
  Text(0.25, 0.1, 'gini = 0.337\nsamples = 21\nvalue = [22, 6]\nclass = Yes'),
  Text(0.39285714285714285, 0.5, 'BMI <= 45.1\ngini = 0.388\nsamples = 23\nvalue
= [10, 28]\nclass = No'),
  Text(0.35714285714285715, 0.3, 'Pregnancies <= 0.5\ngini = 0.312\nsamples =
18\nvalue = [6, 25]\nclass = No'),
  Text(0.32142857142857145, 0.1, 'gini = 0.0\nsamples = 5\nvalue = [0, 10]\nclass
= No'),
  Text(0.39285714285714285, 0.1, 'gini = 0.408\nsamples = 13\nvalue = [6,
15]\nclass = No'),
  Text(0.42857142857142855, 0.3, 'gini = 0.49\nsamples = 5\nvalue = [4, 3]\nclass
= Yes'),
  Text(0.7142857142857143, 0.7, 'Glucose <= 147.5\ngini = 0.5\nsamples =
182\nvalue = [137, 136]\nclass = Yes'),
  Text(0.5714285714285714, 0.5, 'Glucose <= 94.0\ngini = 0.476\nsamples =
135\nvalue = [126, 81]\nclass = Yes'),
  Text(0.5, 0.3, 'SkinThickness <= 30.5\ngini = 0.188\nsamples = 28\nvalue = [34,
4]\nclass = Yes'),
  Text(0.4642857142857143, 0.1, 'gini = 0.069\nsamples = 21\nvalue = [27,
1]\nclass = Yes'),
  Text(0.5357142857142857, 0.1, 'gini = 0.42\nsamples = 7\nvalue = [7, 3]\nclass
= Yes'),
  Text(0.6428571428571429, 0.3, 'BMI <= 30.2\ngini = 0.496\nsamples = 107\nvalue
= [92, 77]\nclass = Yes'),
  Text(0.6071428571428571, 0.1, 'gini = 0.389\nsamples = 42\nvalue = [50,
18]\nclass = Yes'),
  Text(0.6785714285714286, 0.1, 'gini = 0.486\nsamples = 65\nvalue = [42,
59]\nclass = No'),
  Text(0.8571428571428571, 0.5, 'BloodPressure <= 67.0\ngini = 0.278\nsamples =
47\nvalue = [11, 55]\nclass = No'),
  Text(0.7857142857142857, 0.3, 'DiabetesPedigreeFunction <= 0.268\ngini =
0.1\nsamples = 12\nvalue = [1, 18]\nclass = No'),
  Text(0.75, 0.1, 'gini = 0.245\nsamples = 5\nvalue = [1, 6]\nclass = No'),
  Text(0.8214285714285714, 0.1, 'gini = 0.0\nsamples = 7\nvalue = [0, 12]\nclass
= No'),
  Text(0.9285714285714286, 0.3, 'Age <= 38.0\ngini = 0.335\nsamples = 35\nvalue =
[10, 37]\nclass = No'),
  Text(0.8928571428571429, 0.1, 'gini = 0.465\nsamples = 15\nvalue = [7,
12]\nclass = No'),
  Text(0.9642857142857143, 0.1, 'gini = 0.191\nsamples = 20\nvalue = [3,
25]\nclass = No')]

```

