

6pjgzkdea

August 2, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: df=pd.read_csv("/content/drive/MyDrive/mydatasets/C6_bmi.csv")
df
```

```
[3]:
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
..
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

[500 rows x 4 columns]

```
[4]: df.head()
```

```
[4]:
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3

1 Data Cleaning and Data Preprocessing

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 4 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Gender  500 non-null     object
 1   Height  500 non-null     int64
 2   Weight  500 non-null     int64
 3   Index   500 non-null     int64
dtypes: int64(3), object(1)
memory usage: 15.8+ KB
```

```
[6]: df.describe()
```

```
[6]:
```

	Height	Weight	Index
count	500.000000	500.000000	500.000000
mean	169.944000	106.000000	3.748000
std	16.375261	32.382607	1.355053
min	140.000000	50.000000	0.000000
25%	156.000000	80.000000	3.000000
50%	170.500000	106.000000	4.000000
75%	184.000000	136.000000	5.000000
max	199.000000	160.000000	5.000000

```
[7]: df.columns
```

```
[7]: Index(['Gender', 'Height', 'Weight', 'Index'], dtype='object')
```

```
[8]: feature_matrix = df.iloc[:,1:3]
target_vector = df.iloc[:,-1]
```

```
[9]: fs = StandardScaler().fit_transform(feature_matrix)
logr = LogisticRegression()
logr.fit(fs,target_vector)
```

```
[9]: LogisticRegression()
```

```
[10]: observation=[[1,2]]
prediction = logr.predict(observation)
print(prediction)
```

```
[5]
```

```
[11]: logr.classes_
```

```
[11]: array([0, 1, 2, 3, 4, 5])
```

```
[12]: logr.predict_proba(observation)
```

```
[12]: array([[5.59566976e-11, 6.05990036e-10, 1.19071465e-07, 4.99471797e-05,
          2.03791363e-02, 9.79570797e-01]])
```

Random Forest

```
[22]: df
```

```
[22]:
```

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
..
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

[500 rows x 4 columns]

```
[23]: g1={"Gender":{"Male":1, "Female":2}}
df=df.replace(g1)
df
```

```
[23]:
```

	Gender	Height	Weight	Index
0	1	174	96	4
1	1	189	87	2
2	2	185	110	4
3	2	195	104	3
4	1	149	61	3
..
495	2	150	153	5
496	2	184	121	4
497	2	141	136	5
498	1	150	95	5
499	1	173	131	5

[500 rows x 4 columns]

```
[25]: x=df.drop('Gender', axis=1)
      y=df['Gender']
```

```
[26]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
[27]: from sklearn.ensemble import RandomForestClassifier
      rfc = RandomForestClassifier()
      rfc.fit(x_train,y_train)
```

```
[27]: RandomForestClassifier()
```

```
[28]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators': [10,20,30,40,50]
                    }
```

```
[29]: from sklearn.model_selection import GridSearchCV
      grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
      grid_search.fit(x_train,y_train)
```

```
[29]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [1, 2, 3, 4, 5],
                              'min_samples_leaf': [5, 10, 15, 20, 25],
                              'n_estimators': [10, 20, 30, 40, 50]},
                  scoring='accuracy')
```

```
[30]: grid_search.best_score_
```

```
[30]: 0.5628571428571428
```

```
[31]: rfc_best = grid_search.best_estimator_
```

```
[32]: from sklearn.tree import plot_tree
      plt.figure(figsize=(89,40))
      plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes', 'No'],
                filled=True)
```

```
[32]: [Text(0.5, 0.75, 'Weight <= 67.5\ngini = 0.496\nsamples = 211\nvalue = [159,
191]\nclass = No'),
      Text(0.25, 0.25, 'gini = 0.353\nsamples = 28\nvalue = [11, 37]\nclass = No'),
      Text(0.75, 0.25, 'gini = 0.5\nsamples = 183\nvalue = [148, 154]\nclass = No')]
```

