

airk8vzg6

August 2, 2023

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[4]: df_train=pd.read_csv("/content/drive/MyDrive/mydatasets/C8_loan-train.csv")
df_test=pd.read_csv("/content/drive/MyDrive/mydatasets/C8_loan-test.csv")
```

```
[5]: df_train.dropna(inplace=True)
df_test.dropna(inplace=True)
```

```
[6]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 480 entries, 1 to 613
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Loan_ID               480 non-null   object
 1   Gender                480 non-null   object
 2   Married               480 non-null   object
 3   Dependents            480 non-null   object
 4   Education             480 non-null   object
 5   Self_Employed         480 non-null   object
 6   ApplicantIncome       480 non-null   int64
 7   CoapplicantIncome     480 non-null   float64
 8   LoanAmount            480 non-null   float64
 9   Loan_Amount_Term      480 non-null   float64
10   Credit_History         480 non-null   float64
11   Property_Area         480 non-null   object
12   Loan_Status           480 non-null   object
```

```
dtypes: float64(4), int64(1), object(8)
memory usage: 52.5+ KB
```

```
[7]: df_train.columns
```

```
[7]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
         'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
         'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
        dtype='object')
```

```
[8]: feature_matrix = df_train[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
                                'Loan_Amount_Term', 'Credit_History']]  
target_vector = df_train[['Self_Employed']]
```

```
[9]: fs = StandardScaler().fit_transform(feature_matrix)
logr = LogisticRegression()
logr.fit(fs,target_vector)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    y = column_or_1d(y, warn=True)

```

```
[9]: LogisticRegression()
```

```
[10]: observation = df_test[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
                             'Loan_Amount_Term', 'Credit_History']]
prediction = logr.predict(observation)
print(prediction)
```

[illegible]

```
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes']
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has
feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
[11]: logr.classes_
```

```
[11]: array(['No', 'Yes'], dtype=object)
```

```
[12]: logr.predict_proba(observation)[0][0]
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has
feature names, but LogisticRegression was fitted without feature names
warnings.warn(
```

```
[12]: 0.0
```

Random Forest

```
[22]: df2=pd.read_csv("/content/drive/MyDrive/mydatasets/C8_loan-test.csv")
df2.dropna(inplace=True)
df2['Self_Employed'].value_counts()
```

```
[22]: No      257
      Yes      32
      Name: Self_Employed, dtype: int64
```

```
[23]: x=df2[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
            'Loan_Amount_Term', 'Credit_History']]
      y=df2['Self_Employed']
```

```
[24]: g1={'Self_Employed':{'No':1, 'Yes':2}}
      df2=df2.replace(g1)
      df2
```

```
[24]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001015	Male	Yes	0	Graduate	1	
1	LP001022	Male	Yes	1	Graduate	1	
2	LP001031	Male	Yes	2	Graduate	1	
4	LP001051	Male	No	0	Not Graduate	1	
5	LP001054	Male	Yes	0	Not Graduate	2	

```

..      ...      ...      ...      ...      ...      ...
361 LP002969 Male Yes 1 Graduate 1
362 LP002971 Male Yes 3+ Not Graduate 2
363 LP002975 Male Yes 0 Graduate 1
365 LP002986 Male Yes 0 Graduate 1
366 LP002989 Male No 0 Graduate 2

ApplicantIncome CoapplicantIncome LoanAmount Loan_Amount_Term \
0 5720 0 110.0 360.0
1 3076 1500 126.0 360.0
2 5000 1800 208.0 360.0
4 3276 0 78.0 360.0
5 2165 3422 152.0 360.0
..      ...      ...      ...      ...
361 2269 2167 99.0 360.0
362 4009 1777 113.0 360.0
363 4158 709 115.0 360.0
365 5000 2393 158.0 360.0
366 9200 0 98.0 180.0

Credit_History Property_Area
0 1.0 Urban
1 1.0 Urban
2 1.0 Urban
4 1.0 Urban
5 1.0 Urban
..      ...      ...
361 1.0 Semiurban
362 1.0 Urban
363 1.0 Urban
365 1.0 Rural
366 1.0 Rural

```

[289 rows x 12 columns]

```
[25]: from sklearn.model_selection import train_test_split
```

```
[26]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
[27]: from sklearn.ensemble import RandomForestClassifier
```

```
[28]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
[28]: RandomForestClassifier()
```

```
[29]: parameters={'max_depth': [1,2,3,4,5],
                'min_samples_leaf': [5,10,15,20,25],
                'n_estimators': [10,20,30,40,50]
            }
```

```
[30]: from sklearn.model_selection import GridSearchCV
      grid_search_
      ↪=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
      grid_search.fit(x_train,y_train)
```

```
[30]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                  param_grid={'max_depth': [1, 2, 3, 4, 5],
                              'min_samples_leaf': [5, 10, 15, 20, 25],
                              'n_estimators': [10, 20, 30, 40, 50]},
                  scoring='accuracy')
```

```
[31]: grid_search.best_score_
```

```
[31]: 0.8762376237623762
```

```
[32]: rfc_best=grid_search.best_estimator_
```

```
[33]: from sklearn.tree import plot_tree

      plt.figure(figsize=(80,40))
      plot_tree(rfc_best.estimators_[5],feature_names=x.
      ↪columns,class_names=['a','b'],filled=True)
```

```
[33]: [Text(0.5, 0.75, 'LoanAmount <= 85.0\n'gini = 0.253\n'nsamples = 125\n'nvalue = [172,
30]\n'nclass = a'),
      Text(0.25, 0.25, 'gini = 0.0\n'nsamples = 16\n'nvalue = [24, 0]\n'nclass = a'),
      Text(0.75, 0.25, 'gini = 0.28\n'nsamples = 109\n'nvalue = [148, 30]\n'nclass = a')]
```

